

rocchio-assign2

July 10, 2023

```
[10]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_excel("USStates.xlsx")
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   State                 50 non-null    object
 1   Region               50 non-null    object
 2   Population            50 non-null    float64
 3   HouseholdIncome       50 non-null    float64
 4   HighSchool           50 non-null    float64
 5   College              50 non-null    float64
 6   Smokers               50 non-null    float64
 7   PhysicalActivity      50 non-null    float64
 8   Obese                50 non-null    float64
 9   NonWhite             50 non-null    float64
10   HeavyDrinkers        50 non-null    float64
11   TwoParents           50 non-null    float64
12   Insured              50 non-null    float64
dtypes: float64(11), object(2)
memory usage: 5.2+ KB
```

1 Task 1

- 1.1 Technically, all variables below with the exception of the State could be considered both explanatory and response variables with some being more viable than others.

Column	Considered Explanatory	Considered Response	Note
State	—	—	This would be considered an identification variable.
Region	X	X	This could be both a target variable & an explanatory variable.
Population	X	X	This could be both a target variable & an explanatory variable.
HouseholdIncome	X	X	This could be both a target variable & an explanatory variable.
HighSchool	X	X	This could be both a target variable & an explanatory variable.
College	X	X	This could be both a target variable & an explanatory variable.
Smokers	X	X	This could be both a target variable & an explanatory variable.
PhysicalActivity	X	X	This could be both a target variable & an explanatory variable.
Obese	X	X	This could be both a target variable & an explanatory variable.
NonWhite	X	X	This could be both a target variable & an explanatory variable.
HeavyDrinkers	X	X	This could be both a target variable & an explanatory variable.
TwoParents	X	X	This could be both a target variable & an explanatory variable.
Insured	X	X	This could be both a target variable & an explanatory variable.

2 Task 2

- 2.1 The population of interest in this dataset would be residents of the 50 states represented, assuming these data points are representative of the state as a whole. If the dataset is a representative sample, it is reasonable to generalize the findings to the larger population in each of these states. Given that each data point is labeled with 'State', the population of interest is likely the entire U.S. population as subdivided by state. It's essential to remember that this dataset might not perfectly capture the entire population of each state or the U.S., but it is being used to make generalized assumptions about those populations.

3 Task 3

```
[11]: summary = df.describe(include='all')
summary = summary.loc[['count', 'mean', 'std']]
pd.DataFrame(summary.T)
```

```
[11]:
```

	count	mean	std
State	50	NaN	NaN
Region	50	NaN	NaN
Population	50.0	6.36394	7.15096
HouseholdIncome	50.0	53.28428	8.690234
HighSchool	50.0	89.32	3.107135
College	50.0	30.83	6.078643
Smokers	50.0	19.316	3.523122
PhysicalActivity	50.0	50.734	5.509643
Obese	50.0	28.766	3.369286
NonWhite	50.0	22.156	12.685572
HeavyDrinkers	50.0	6.046	1.175292
TwoParents	50.0	65.524	5.17074
Insured	50.0	80.148	5.494087

```
[12]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

explanatory_vars = ['HighSchool', 'College', 'Smokers', 'PhysicalActivity',
                    'Obese', 'NonWhite', 'HeavyDrinkers', 'TwoParents',
                    ↪ 'Insured']

fig, axes = plt.subplots(nrows=2, ncols=5, figsize=(25, 15))

for i, var in enumerate(explanatory_vars):
    row = i // 5
    col = i % 5
```

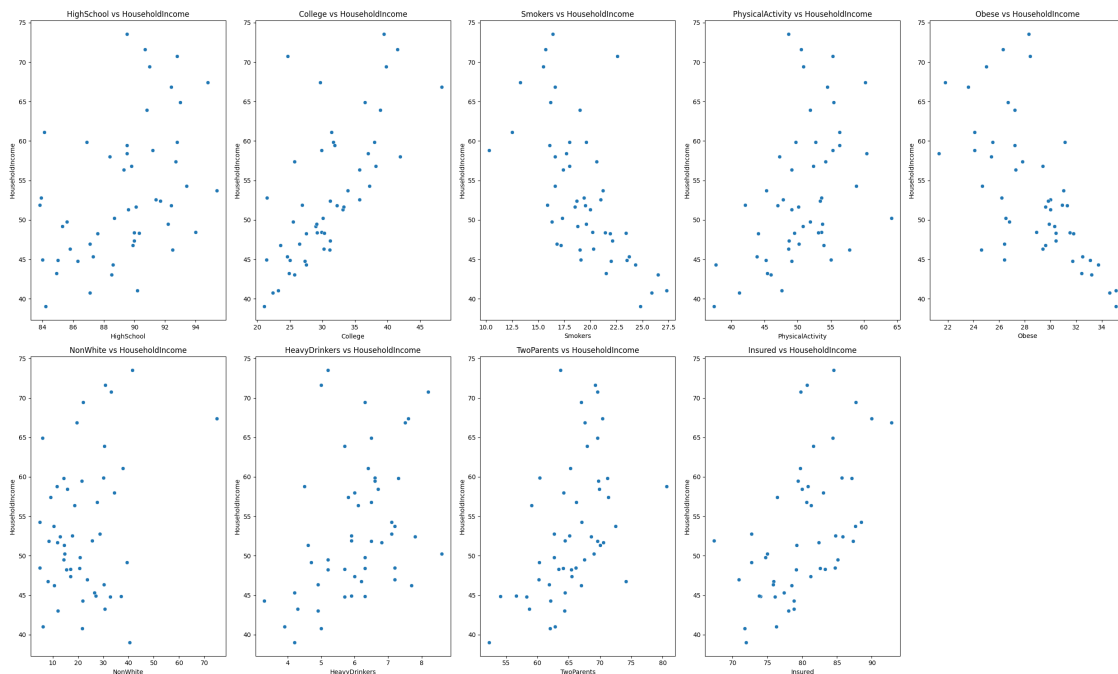
```

ax = axes[row, col]
sns.scatterplot(x=df[var], y=df['HouseholdIncome'], ax=ax)
ax.set_title(f'{var} vs HouseholdIncome')
ax.set_xlabel(var)
ax.set_ylabel('HouseholdIncome')

if len(explanatory_vars) < 10:
    for i in range(len(explanatory_vars), 10):
        row = i // 5
        col = i % 5
        fig.delaxes(axes[row][col])

plt.tight_layout()
plt.show()

```



4 Task 4

```

[13]: correlation = df[explanatory_vars+['HouseholdIncome']].
      ↪corr()['HouseholdIncome'].drop('HouseholdIncome')
      print(correlation)

```

HighSchool	0.430845
College	0.685591
Smokers	-0.637522
PhysicalActivity	0.440417

```

Obese          -0.649112
NonWhite       0.252942
HeavyDrinkers  0.373014
TwoParents     0.477644
Insured        0.549679
Name: HouseholdIncome, dtype: float64

```

```

[14]: for var in explanatory_vars:
        model = sm.OLS(df['HouseholdIncome'], sm.add_constant(df[var])).fit()

        print(f"Regression model for {var}:")
        ↵
        ↪print('-----')
        print('')
        print(model.summary())

        # Residuals vs Fitted values
        plt.figure(figsize=(8,4))
        plt.subplot(121)
        plt.scatter(df[var], model.resid)
        plt.axhline(y=0, color='r', linestyle='-')
        plt.xlabel(var)
        plt.ylabel('Residuals')
        plt.title(f"Residuals vs {var}")

        # Normal Q-Q plot
        plt.subplot(122)
        sm.qqplot(model.resid, line='q', ax=plt.gca())
        plt.title(Normal Q-Q plot)

        plt.tight_layout()
        plt.show()
        print('\n\n')
        ↵
        ↪print('#####')
        ↵
        ↪print('#####')
        ↵
        ↪print('#####')
        print('\n\n')

```

Regression model for HighSchool:

```

-----

                        OLS Regression Results
=====
Dep. Variable:          HouseholdIncome   R-squared:                0.186

```

```

Model: OLS Adj. R-squared: 0.169
Method: Least Squares F-statistic: 10.94
Date: Mon, 10 Jul 2023 Prob (F-statistic): 0.00179
Time: 15:24:42 Log-Likelihood: -173.42
No. Observations: 50 AIC: 350.8
Df Residuals: 48 BIC: 354.7
Df Model: 1
Covariance Type: nonrobust

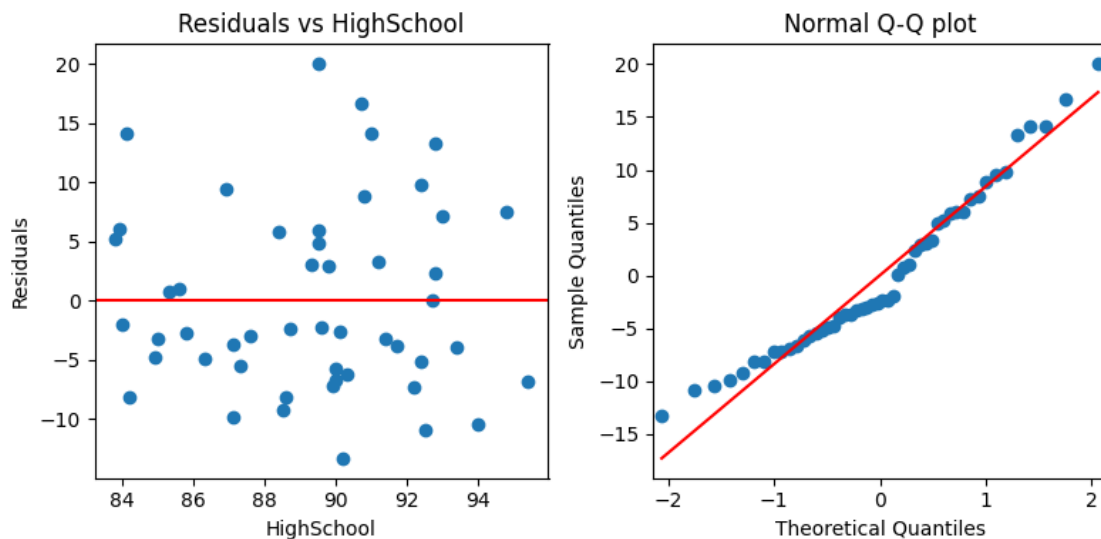
```

	coef	std err	t	P> t	[0.025	0.975]
const	-54.3476	32.559	-1.669	0.102	-119.811	11.116
HighSchool	1.2050	0.364	3.308	0.002	0.473	1.937

Omnibus:	3.700	Durbin-Watson:	1.761
Prob(Omnibus):	0.157	Jarque-Bera (JB):	3.542
Skew:	0.629	Prob(JB):	0.170
Kurtosis:	2.660	Cond. No.	2.60e+03

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.6e+03. This might indicate that there are strong multicollinearity or other numerical problems.



```
#####
#####
#####
```

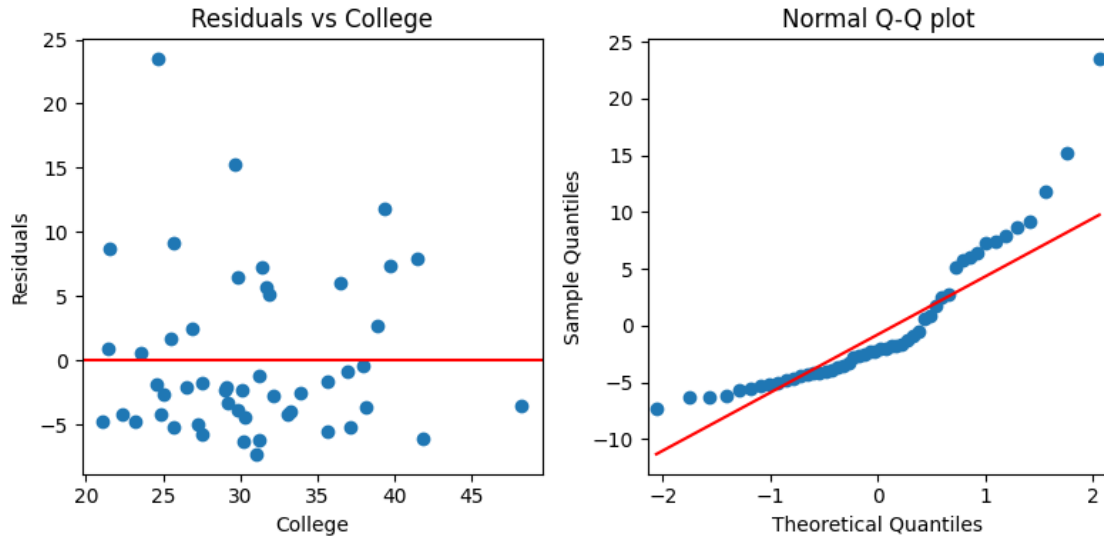
Regression model for College:

OLS Regression Results						
=====						
Dep. Variable:	HouseholdIncome	R-squared:	0.470			
Model:	OLS	Adj. R-squared:	0.459			
Method:	Least Squares	F-statistic:	42.57			
Date:	Mon, 10 Jul 2023	Prob (F-statistic):	3.94e-08			
Time:	15:24:42	Log-Likelihood:	-162.68			
No. Observations:	50	AIC:	329.4			
Df Residuals:	48	BIC:	333.2			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	23.0664	4.719	4.888	0.000	13.579	32.554
College	0.9801	0.150	6.525	0.000	0.678	1.282
=====						
Omnibus:	22.384	Durbin-Watson:	1.811			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	32.343			
Skew:	1.539	Prob(JB):	9.48e-08			
Kurtosis:	5.461	Cond. No.	164.			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
#####
#####
#####
```

Regression model for Smokers:

OLS Regression Results						
=====						
Dep. Variable:	HouseholdIncome		R-squared:		0.406	
Model:	OLS		Adj. R-squared:		0.394	
Method:	Least Squares		F-statistic:		32.87	
Date:	Mon, 10 Jul 2023		Prob (F-statistic):		6.40e-07	
Time:	15:24:42		Log-Likelihood:		-165.51	
No. Observations:	50		AIC:		335.0	
Df Residuals:	48		BIC:		338.8	
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

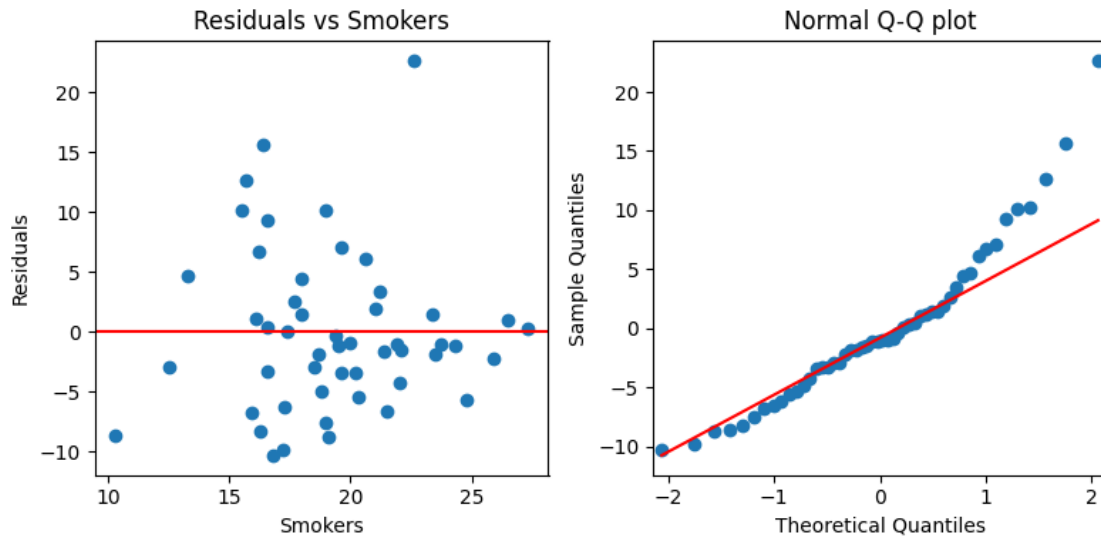
const	83.6593	5.384	15.539	0.000	72.834	94.484
Smokers	-1.5725	0.274	-5.733	0.000	-2.124	-1.021
=====						

Omnibus:	13.202	Durbin-Watson:	2.153
Prob(Omnibus):	0.001	Jarque-Bera (JB):	14.208
Skew:	1.073	Prob(JB):	0.000822
Kurtosis:	4.488	Cond. No.	111.

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



 #####
 #####

Regression model for PhysicalActivity:

OLS Regression Results			
=====			
Dep. Variable:	HouseholdIncome	R-squared:	0.194
Model:	OLS	Adj. R-squared:	0.177
Method:	Least Squares	F-statistic:	11.55
Date:	Mon, 10 Jul 2023	Prob (F-statistic):	0.00137
Time:	15:24:42	Log-Likelihood:	-173.16

```

No. Observations:      50    AIC:      350.3
Df Residuals:          48    BIC:      354.1
Df Model:              1
Covariance Type:      nonrobust
=====

```

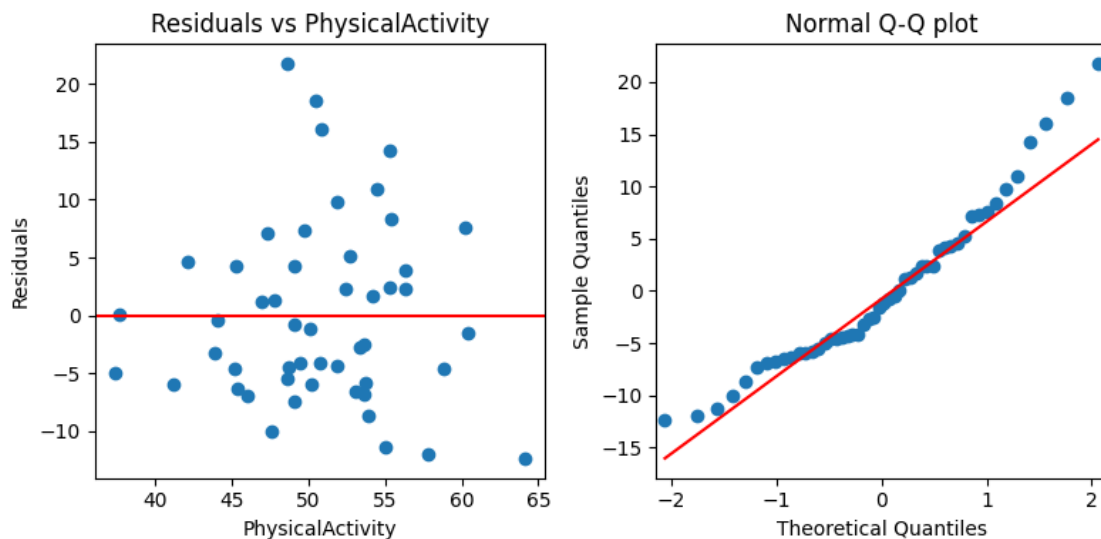
```

=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const          18.0414      10.429       1.730     0.090     -2.928
39.011
PhysicalActivity  0.6947       0.204       3.399     0.001      0.284
1.106
=====
Omnibus:          6.248    Durbin-Watson:      2.177
Prob(Omnibus):    0.044    Jarque-Bera (JB):    5.484
Skew:             0.799    Prob(JB):            0.0644
Kurtosis:         3.277    Cond. No.            478.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



#####

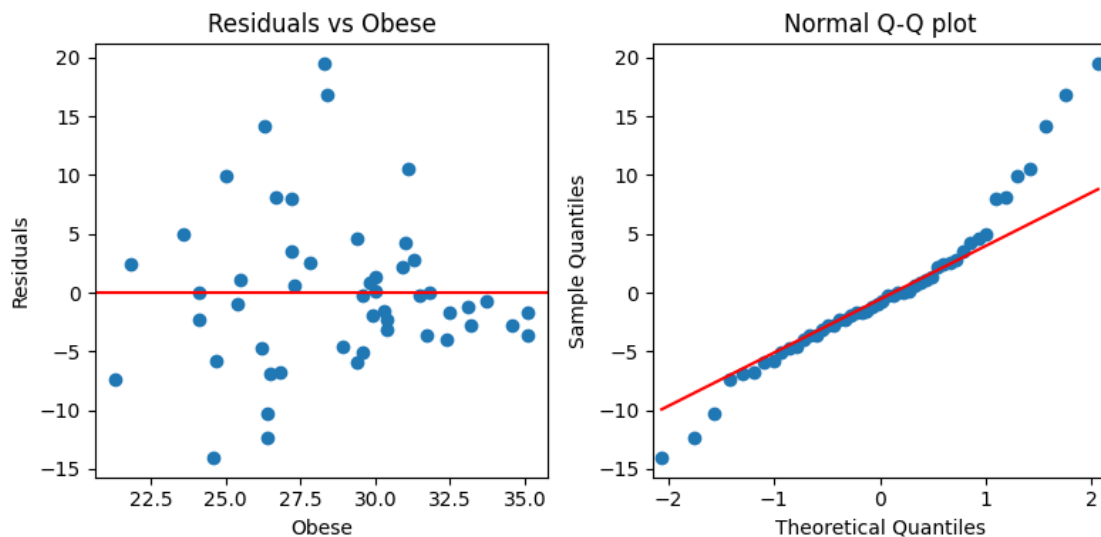
#####

Regression model for Obese:

```
-----
                        OLS Regression Results
=====
Dep. Variable:      HouseholdIncome    R-squared:                0.421
Model:              OLS                Adj. R-squared:           0.409
Method:             Least Squares      F-statistic:             34.95
Date:               Mon, 10 Jul 2023    Prob (F-statistic):      3.42e-07
Time:               15:24:42           Log-Likelihood:          -164.88
No. Observations:   50                AIC:                    333.8
Df Residuals:       48                BIC:                    337.6
Df Model:           1
Covariance Type:    nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const         101.4449      8.201     12.370     0.000     84.956    117.934
Obese          -1.6742      0.283     -5.912     0.000     -2.244     -1.105
=====
Omnibus:                8.503    Durbin-Watson:           2.282
Prob(Omnibus):          0.014    Jarque-Bera (JB):         7.865
Skew:                   0.775    Prob(JB):                 0.0196
Kurtosis:               4.171    Cond. No.                  252.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
#####
#####
#####
```

Regression model for NonWhite:

OLS Regression Results						
=====						
Dep. Variable:	HouseholdIncome		R-squared:		0.064	
Model:	OLS		Adj. R-squared:		0.044	
Method:	Least Squares		F-statistic:		3.281	
Date:	Mon, 10 Jul 2023		Prob (F-statistic):		0.0763	
Time:	15:24:42		Log-Likelihood:		-176.90	
No. Observations:	50		AIC:		357.8	
Df Residuals:	48		BIC:		361.6	
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

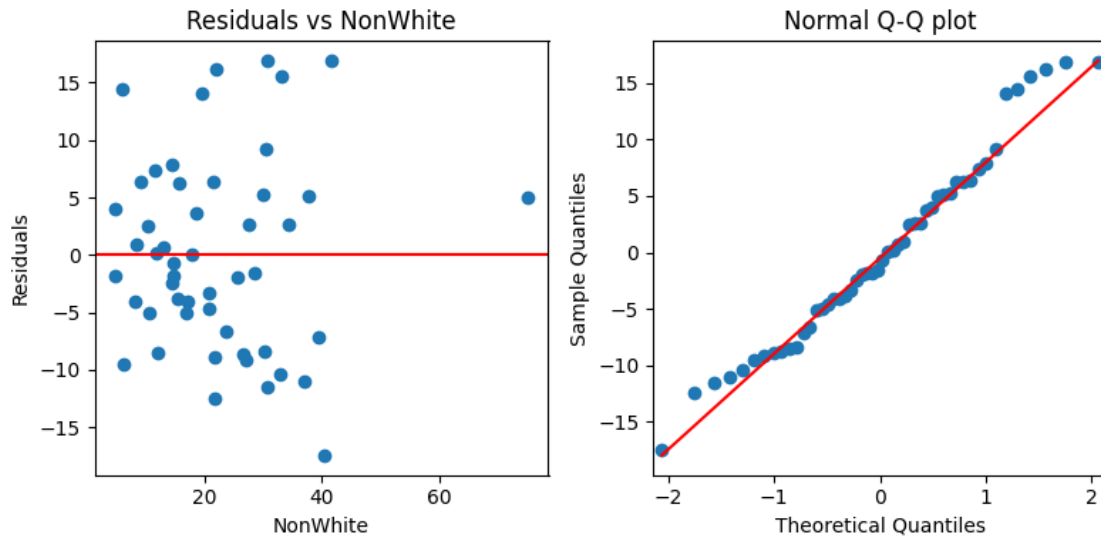
const	49.4451	2.436	20.295	0.000	44.547	54.344
NonWhite	0.1733	0.096	1.811	0.076	-0.019	0.366
=====						

Omnibus:	1.522	Durbin-Watson:	1.830
Prob(Omnibus):	0.467	Jarque-Bera (JB):	1.491
Skew:	0.337	Prob(JB):	0.475
Kurtosis:	2.490	Cond. No.	51.7

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



 #####
 #####

Regression model for HeavyDrinkers:

OLS Regression Results			
=====			
Dep. Variable:	HouseholdIncome	R-squared:	0.139
Model:	OLS	Adj. R-squared:	0.121
Method:	Least Squares	F-statistic:	7.758
Date:	Mon, 10 Jul 2023	Prob (F-statistic):	0.00763
Time:	15:24:42	Log-Likelihood:	-174.81

No. Observations: 50 AIC: 353.6
Df Residuals: 48 BIC: 357.4
Df Model: 1
Covariance Type: nonrobust

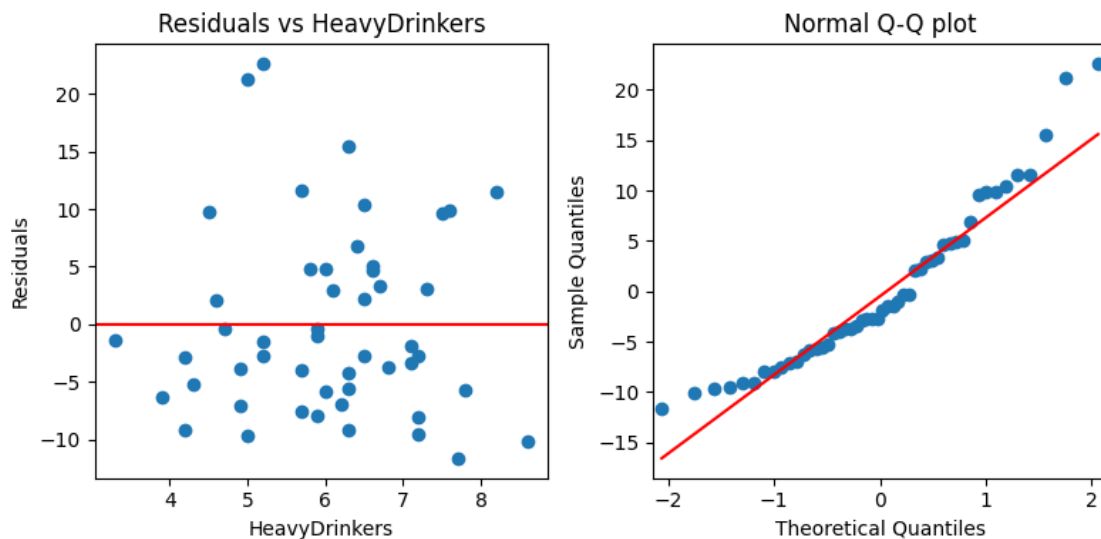
	coef	std err	t	P> t	[0.025
0.975]					

-					
const	36.6088	6.097	6.005	0.000	24.350
48.867					
HeavyDrinkers	2.7581	0.990	2.785	0.008	0.767
4.749					

Omnibus:	7.882	Durbin-Watson:	1.912		
Prob(Omnibus):	0.019	Jarque-Bera (JB):	7.278		
Skew:	0.917	Prob(JB):	0.0263		
Kurtosis:	3.361	Cond. No.	33.4		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



#####

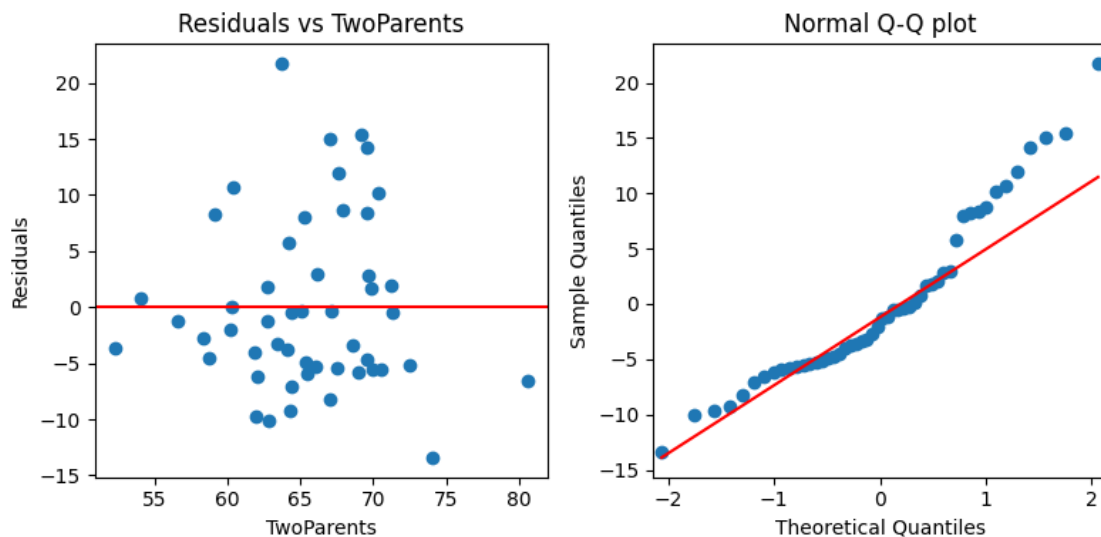
#####

Regression model for TwoParents:

```
-----
                        OLS Regression Results
=====
Dep. Variable:          HouseholdIncome    R-squared:                0.228
Model:                  OLS                Adj. R-squared:          0.212
Method:                 Least Squares      F-statistic:             14.19
Date:                   Mon, 10 Jul 2023    Prob (F-statistic):      0.000452
Time:                   15:24:43           Log-Likelihood:          -172.08
No. Observations:       50                AIC:                     348.2
Df Residuals:           48                BIC:                     352.0
Df Model:               1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          0.6845       14.007        0.049      0.961     -27.479     28.848
TwoParents     0.8028        0.213        3.767      0.000         0.374         1.231
=====
Omnibus:                 6.567    Durbin-Watson:           1.729
Prob(Omnibus):           0.038    Jarque-Bera (JB):         6.057
Skew:                    0.850    Prob(JB):                 0.0484
Kurtosis:                3.143    Cond. No.                  844.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



 #####
 #####

Regression model for Insured:

=====						
OLS Regression Results						
=====						
Dep. Variable:	HouseholdIncome	R-squared:	0.302			
Model:	OLS	Adj. R-squared:	0.288			
Method:	Least Squares	F-statistic:	20.78			
Date:	Mon, 10 Jul 2023	Prob (F-statistic):	3.56e-05			
Time:	15:24:43	Log-Likelihood:	-169.56			
No. Observations:	50	AIC:	343.1			
Df Residuals:	48	BIC:	346.9			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-16.4004	15.321	-1.070	0.290	-47.205	14.405
Insured	0.8695	0.191	4.559	0.000	0.486	1.253
=====						

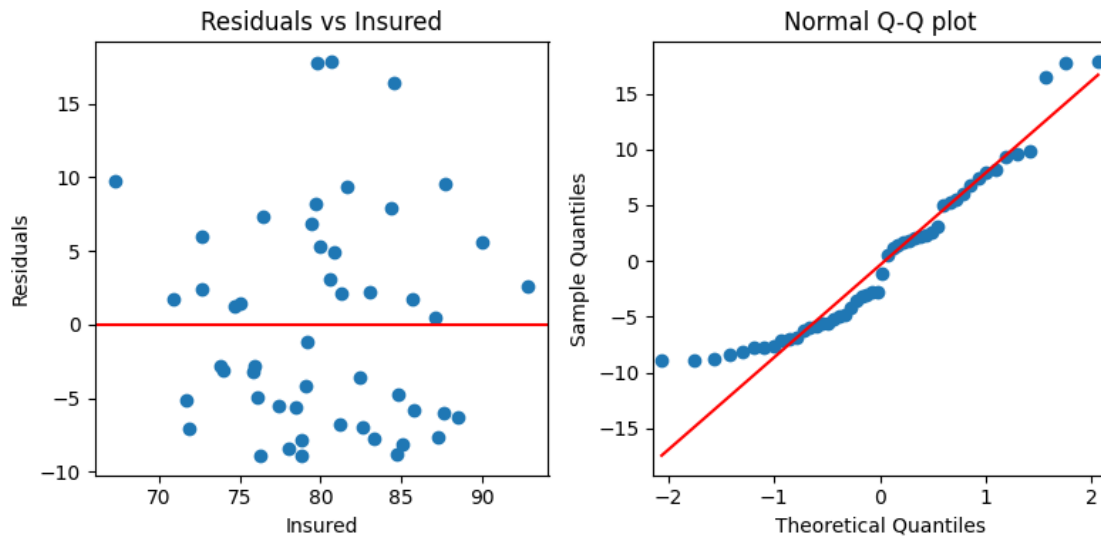
Omnibus:	4.875	Durbin-Watson:	1.820
Prob(Omnibus):	0.087	Jarque-Bera (JB):	4.670
Skew:	0.741	Prob(JB):	0.0968
Kurtosis:	2.791	Cond. No.	1.19e+03

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.19e+03. This might indicate that there are strong multicollinearity or other numerical problems.



 #####
 #####

4.1 Note on simple linear regression: The value of the modelling method would depend on the following factors:

4.1.1 Linearity:

The assumption of linearity is that the relationship between the independent and dependent variables should ideally be represented by a straight line. If the scatter plots we created during step 3 indicate a straight-line trend between these variables, then this requirement is met.

4.1.2 Normality:

Another assumption we make is that the residuals follow a normal distribution. We can verify this using the Omnibus test, where a p-value of less than 0.05 points to non-normality. Here, the Omnibus test p-values for the 'Obese', 'HeavyDrinkers', and 'TwoParents' models are all under 0.05, suggesting possible breaches of the normal curve assumption.

4.1.3 Homoscedasticity:

Homoscedasticity suggests that the variance of the residuals should remain constant at all levels of the independent variables. If our scatter plots from step 3 are showing a funnel shape, it might imply that we've strayed from this principle.

4.1.4 Significant R-squared values:

The R-squared value tells us how much of the variance in the dependent variable is explained by the independent one. For example above, Obese and Insured have higher R-squared values (0.421 and 0.302), showing that they could be strong predictors of HouseholdIncome. For most other variables though, the R-squared values are quite low, meaning that simple linear regression might not be the most effective approach.

4.1.5 Lack of multicollinearity:

With these simple linear regression models, this isn't a problem because we are using a single input variable.

5 Task 5

```
[15]: import statsmodels.formula.api as smf

model1 = smf.ols(formula='HouseholdIncome ~ College', data=df)
results1 = model1.fit()
print(results1.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          HouseholdIncome    R-squared:                0.470
Model:                  OLS               Adj. R-squared:         0.459
Method:                 Least Squares      F-statistic:             42.57
Date:                   Mon, 10 Jul 2023    Prob (F-statistic):       3.94e-08
Time:                   15:24:43           Log-Likelihood:          -162.68
No. Observations:       50                AIC:                   329.4
Df Residuals:           48                BIC:                   333.2
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept              23.0664      4.719      4.888      0.000      13.579      32.554
```

College	0.9801	0.150	6.525	0.000	0.678	1.282
=====						
Omnibus:		22.384	Durbin-Watson:			1.811
Prob(Omnibus):		0.000	Jarque-Bera (JB):			32.343
Skew:		1.539	Prob(JB):			9.48e-08
Kurtosis:		5.461	Cond. No.			164.
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

5.1 Analysis of College Variable

The College Education variable is likely highly correlated to earnings because it's widely known that education often plays a significant role in income levels. People with higher degrees, like a college diploma, usually earn more, leading to a higher total household income. The prediction equation for Model 1 would be:

$$\text{HouseholdIncome} = 23.0664 + 0.9801 * \text{College}$$

Here's how to interpret the coefficients:

5.1.1 Constant (Intercept):

This value is the estimated average income for households where the College factor is zero, meaning no college education. So, if we're talking about someone who hasn't been to college, their expected average household income would be roughly \$23,066.4.

5.1.2 College = 0.9801:

This number is the slope of the line and represents the expected change in household income for each unit increase in 'College'. So, whether that means one additional year of college education or making the jump from no degree to having one (depending on how we define 'College'), we can expect to see an approximate \$980.1 increase in household income, all other things being equal.

5.1.3 The R-squared statistic for Model 1 is 0.470:

This metric tells us what proportion of the variance in household income can be predicted from the 'College' factor. This means that about 47% of the changes in income are explained by someone's education level. That is a moderate R-squared value, and while higher values might sound better, they can also mean the model is overfitting the data.

5.1.4 As for why we might start with the College variable as an explanatory variable, it might be due to an initial hypothesis that education, and in particular college education, has a strong impact on household income. It's a reasonable first variable to consider given common knowledge and previous research showing that education level is often correlated with income. However, it's always necessary to conduct the actual analysis to confirm whether the data supports this hypothesis in any specific instance.

6 Task 6

```
[16]: anova_table = sm.stats.anova_lm(results1, typ=2)
      print(anova_table)
```

	sum_sq	df	F	PR(>F)
College	1739.358780	1.0	42.572008	3.941054e-08
Residual	1961.129512	48.0	NaN	NaN

6.1 From the OLS Regression results and the ANOVA table, we can construct our hypotheses and interpretations as follows:

6.1.1 For the parameter associated with the 'College' explanatory variable:

Null Hypothesis (H0): The coefficient of the College variable (Beta_College) equals 0. Alternative Hypothesis (H1): The coefficient of the College variable (Beta_College) is not equal to 0. The T-test result for this hypothesis ($t=6.525$, $P>|t|=0.000$) means that we can reject the null hypothesis. So the College variable has a significant effect on the HouseholdIncome.

6.1.2 For the Intercept parameter:

Null Hypothesis (H0): The Intercept (Beta_Intercept) is equal to 0. Alternative Hypothesis (H1): The Intercept (Beta_Intercept) is not equal to 0. The T-test result for this hypothesis ($t=4.888$, $P>|t|=0.000$) also provides strong evidence to reject the null hypothesis in favor of the alternative. This implies that even when the 'College' variable is equal to zero, the 'HouseholdIncome' is not zero.

6.1.3 From the ANOVA table:

The F statistic value is 42.57 with a p-value (PR(>F)) of 3.94e-08 which is close to zero. This means that we reject the null hypothesis.

6.1.4 In summary, both the 'Intercept' and 'College' parameters are statistically significant predictors of 'HouseholdIncome' and the overall model is significant in predicting the 'HouseholdIncome'.

7 Task 7

```
[1]: library(tidyverse)
```

Attaching core tidyverse packages
2.0.0

tidyverse

```

dplyr      1.1.2      readr      2.1.4
forcats    1.0.0      stringr   1.5.0
ggplot2    3.4.2      tibble     3.2.1
lubridate  1.9.2      tidyr      1.3.0
purrr      1.0.1

```

Conflicts

```
tidyverse_conflicts()
```

```
dplyr::filter() masks stats::filter()
```

```
dplyr::lag() masks stats::lag()
```

Use the conflicted package

(<http://conflicted.r-lib.org/>) to force all conflicts to become errors

```
[2]: coefficients <- read.csv(coefficients.csv)
df <- read.csv("USStates.csv")
model1 <- lm(HouseholdIncome ~ College, data = df)
summary(model1)
```

Call:

```
lm(formula = HouseholdIncome ~ College, data = df)
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-7.319 -4.245 -2.203  2.652 23.484

```

Coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
(Intercept)  23.0664      4.7187   4.888 1.18e-05 ***
College       0.9801      0.1502   6.525 3.94e-08 ***
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.392 on 48 degrees of freedom

Multiple R-squared: 0.47, Adjusted R-squared: 0.459

F-statistic: 42.57 on 1 and 48 DF, p-value: 3.941e-08

```
[3]: prediction <- predict(model1, df)
residuals <- df$HouseholdIncome - prediction
sse <- sum(residuals^2)
mean_Y <- mean(df$HouseholdIncome)
deviations_Y <- df$HouseholdIncome - mean_Y
sst <- sum((df$HouseholdIncome - mean_Y)^2)
ssr <- sum((prediction - mean_Y)^2)
r_squared <- ssr / sst
cat("Sum of squared residuals (SSE):", sse, "\n")
```

```
cat("Sum of squares total (SST):", sst, "\n")
cat("Sum of squares due to regression (SSR):", ssr, "\n")
cat("R-squared:", r_squared, "\n")
```

Sum of squared residuals (SSE): 1961.13
Sum of squares total (SST): 3700.488
Sum of squares due to regression (SSR): 1739.359
R-squared: 0.4700349

7.1 Upon comparing the results obtained from the code with the output provided, we can see that the computed statistics match:

7.1.1 Sum of squared residuals (SSE): 1961.13

7.1.2 Sum of squares total (SST): 3700.488

7.1.3 Sum of squares due to regression (SSR): 1739.359

7.1.4 R-squared: 0.4700349

7.2 These values align with the ANOVA table and R-squared values reported in the regression output.

7.3 The ANOVA table shows the F-statistic. The F-statistic tests the null hypothesis that for all regression coefficients, at least one of them is non-zero. In this case, the F-statistic is 42.57, with a very low p-value (3.94e-08), indicating strong evidence against the null hypothesis.

7.4 The R-squared value of 0.4700349 represents the proportion of the total variation in the response variable (HouseholdIncome) that is explained by the regression model. It indicates that approximately 47 % of the variability in HouseholdIncome can be accounted for by the College variable in the model.

7.5 The computed statistics match the ANOVA table and R-squared value reported in the regression output, confirming the accuracy of the computations. Any minor differences can be attributed to rounding errors or slight variations in the implementation of the statistical functions across software packages. Overall, the similarity between the computed values and the regression output demonstrates the consistency and validity of the analysis.

8 Task 8

```
[4]: # Calculate the residuals for Model 1
residuals_model1 <- df$HouseholdIncome - predict(model1)

# Standardize the residuals
standardized_residuals <- residuals_model1 / sd(residuals_model1)

# Plot histogram of standardized residuals
```

```

hist(standardized_residuals, breaks = 10, col = "lightblue",
     main = "Histogram of Standardized Residuals for Model 1")


# Plot scatterplot of standardized residuals against predicted values
plot(predict(model1), standardized_residuals, pch = 16, col = "darkblue",
     xlab = "Predicted Values (Y_hat)", ylab = "Standardized Residuals",
     main = "Scatterplot of Standardized Residuals vs. Predicted Values for
     ↪Model 1")

# Add a horizontal line at y = 0
abline(h = 0, col = "red", lty = 2)


# Add a lowess smoother to the scatterplot
lowess_line <- lowess(predict(model1), standardized_residuals, f = 2 / 3, iter_
     ↪= 0)
lines(lowess_line, col = "orange")

```

Histogram of Standardized Residuals for Model 1



Scatterplot of Standardized Residuals vs. Predicted Values for Model 1



- 8.1 In the histogram of standardized residuals for Model 1, it shows that the distribution is approximately bell-shaped with a mean of 0 and a standard deviation of 1. The skewness of 1.539 indicates a mid- positive skew, meaning the residuals are slightly skewed to the right. The kurtosis value of 2.461 shows that the distribution has slightly denser tails.
- 8.2 In the scatterplot of standardized residuals vs. predicted values for Model 1, we can see that there is no linear relationship between the standardized residuals and the predicted values. The correlation value of 0 proves this.
- 8.3 Overall, these graphs provide insights into the distribution and relationship of the standardized residuals in Model 1. The histogram gives us an idea about the normality of the residuals, while the scatterplot helps show the linearity and patterns in the residuals.

9 Task 9

```
[5]: model2 <- lm(HouseholdIncome ~ HighSchool, data = df)
summary(model2)
prediction <- predict(model2, df)
residuals <- df$HouseholdIncome - prediction
sse <- sum(residuals^2)
mean_Y <- mean(df$HouseholdIncome)
deviations_Y <- df$HouseholdIncome - mean_Y
sst <- sum((df$HouseholdIncome - mean_Y)^2)
ssr <- sum((prediction - mean_Y)^2)
r_squared <- ssr / sst
cat("Sum of squared residuals (SSE):", sse, "\n")
cat("Sum of squares total (SST):", sst, "\n")
cat("Sum of squares due to regression (SSR):", ssr, "\n")
cat("R-squared:", r_squared, "\n")
```

Call:

```
lm(formula = HouseholdIncome ~ HighSchool, data = df)
```

Residuals:

Min	1Q	Median	3Q	Max
-13.302	-5.647	-2.430	5.687	20.037

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-54.3476	32.5588	-1.669	0.10159
HighSchool	1.2050	0.3643	3.308	0.00179 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.924 on 48 degrees of freedom

Multiple R-squared: 0.1856, Adjusted R-squared: 0.1687
F-statistic: 10.94 on 1 and 48 DF, p-value: 0.001788

Sum of squared residuals (SSE): 3013.577
Sum of squares total (SST): 3700.488
Sum of squares due to regression (SSR): 686.9114
R-squared: 0.1856272

9.1 The analysis of Model 2 is below:

9.1.1 Coefficients:

The intercept coefficient is -54.3476, indicating that when the HighSchool value is zero, the estimated HouseholdIncome is -54.35. However, this coefficient is not statistically significant at the conventional significance level (p-value: 0.10159).

The coefficient for HighSchool is 1.2050, indicating that for a one-unit increase in HighSchool, the estimated HouseholdIncome increases by 1.2050. This coefficient is somewhat significant (p-value: 0.00179).

9.1.2 R-squared and Adjusted R-squared:

The R-squared is 0.1856. This means that about 18.56 % of the variability in HouseholdIncome can be explained by HighSchool. The adj. R-squared value is 0.1687, which take into account the number of predictors and degrees of freedom.

9.1.3 F-statistic and p-value:

The F-statistic is 10.94 with a p-value of 0.001788. This indicates that there is evidence of a significant relationship between the HighSchool variable and HouseholdIncome.

9.1.4 Comparison of Model 1 and Model 2:

The R-squared of model 1 is higher compared to Model 2.

Considering the higher R-squared value and the significance of the coefficients in Model 1, we can conclude that Model 1 is a better model for predicting HouseholdIncome compared to Model 2.

10 Task 10

```
[6]: model3 <- lm(HouseholdIncome ~ Insured, data = df)
      summary(model3)
      prediction <- predict(model3, df)
      residuals <- df$HouseholdIncome - prediction
      sse <- sum(residuals^2)
      mean_Y <- mean(df$HouseholdIncome)
      deviations_Y <- df$HouseholdIncome - mean_Y
```

```
sst <- sum((df$HouseholdIncome - mean_Y)^2)
ssr <- sum((prediction - mean_Y)^2)
r_squared <- ssr / sst
cat("Sum of squared residuals (SSE):", sse, "\n")
cat("Sum of squares total (SST):", sst, "\n")
cat("Sum of squares due to regression (SSR):", ssr, "\n")
cat("R-squared:", r_squared, "\n")
```

Call:

```
lm(formula = HouseholdIncome ~ Insured, data = df)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.896	-5.963	-1.976	5.200	17.865

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-16.4004	15.3210	-1.070	0.29
Insured	0.8695	0.1907	4.559	3.56e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.335 on 48 degrees of freedom

Multiple R-squared: 0.3021, Adjusted R-squared: 0.2876

F-statistic: 20.78 on 1 and 48 DF, p-value: 3.557e-05

Sum of squared residuals (SSE): 2582.398

Sum of squares total (SST): 3700.488

Sum of squares due to regression (SSR): 1118.09

R-squared: 0.3021466

10.1 Model 3 Analysis:

10.1.1 Interpretation of Coefficients:

The y-intercept coefficient is -16.4004 so if the Insured value were to be zero the estimate for the HouseholdIncome would be -16.4004. However, with a p-value of 0.29 this coefficient is not significant. The coefficient for the Insured variable is 0.8695. This signifies that with each unit increase in Insured, we can expect an estimated increase of 0.8695 in the HouseholdIncome, a relationship that is statistically significant.

10.1.2 R-squared and Adjusted R-squared Interpretation:

The R-squared value of 0.3021 reveals that roughly 30.21% of the variations in HouseholdIncome can be accounted for by the Insured variable. As for the adjusted R-squared value for Insured, it comes to 0.2876.

10.1.3 Interpreting the F-statistic and p-value:

With an F-statistic of 20.78, we have significant evidence of a meaningful correlation between the Insured variable and HouseholdIncome.

10.1.4 Model Comparison:

Comparing Model 3 with both Model 1 and Model 2, Model 3's r-squared value of 0.3021 performs worse when compared to Model 1 with an r-squared value of 0.47 but has a higher r-squared Model 2 with 0.1856.

10.1.5 Based on all the information above, the first model that was trained on education is the best predictor for income.